

Ramakrushna Dash

Senior Data Engineer

ramakrushna.work@gmail.com +91 9776955226 linkedin.com/in/ramakrushna-dash-de

PROFESSIONAL SUMMARY

Senior Data Engineer with ~10 years of experience architecting cloud-native data platforms across Insurance and Financial Services domains. Specializes in real-time Kafka streaming pipelines, lakehouse architecture on Databricks and Snowflake, and enterprise-grade PII tokenization and encryption. Proven track record modernizing legacy Informatica ETL ecosystems to scalable, event-driven solutions on Azure and AWS.

WORK EXPERIENCE

Chubb — Insurance Domain

Nov 2024 – Present

Senior Data Engineer

Modern Streaming Platform (Mar 2025 – Present)

- Architected shared hybrid ingestion: on-prem Linux server bridges Kafka and source systems to ADLS Gen2, enabling secure cloud-side Databricks processing — applied consistently across all 3 production pipelines.
- Engineered ChubbIO pipeline — Thailand Auto Rating claim events via on-prem Kafka → Linux bridge → ADLS Gen2 → Databricks medallion (bronze→silver→gold) → Snowflake; 2,000 files per 6-hour batch with automated reconciliation.
- Built DuckCreek policy pipeline — hourly Kafka batch; API call per ASYNC ID retrieves full XML policy documents; AES-256 encryption on-prem (Azure Key Vault) before ADLS transit; decrypted and processed entirely in Databricks memory with zero PII written to persistent storage; PII tokenized via isolated on-prem server; dual reconciliation on files and IDs; 10,000+ policies warehoused.
- Delivered SQL Server → Snowflake migration — 5M+ rows, 25GB including embedded XML columns requiring custom parsing; AES-256 PII tokenization applied to all fields before load.
- Implemented Chubb-standard AES-256 encryption (Azure Key Vault) and field-level PII tokenization with architectural isolation — Databricks has zero network access to on-prem tokenization server by design.
- Managed CI/CD deployment and environment promotion across all pipelines using Databricks Asset Bundles (DAB).

UW Reporting (Nov 2024 – Mar 2025)

- Built data marts integrating Mainframe systems, flat files, and JSON sources using Informatica Cloud (IICS) and Snowflake.
- Delivered foundational data layer consumed by Power BI dashboards supporting underwriting business decisions and KPI reporting.
- Automated ETL pipelines using Python scripting, improving pipeline reliability and reducing manual intervention.

LTIMindtree

Dec 2022 – Oct 2024

Specialist Data Engineer

- Led end-to-end migration of 50+ Informatica PowerCenter workflows to AWS Glue for a large Insurance client, modernizing legacy ETL infrastructure to cloud-native architecture.
- Implemented parallel load reconciliation — running dual pipelines across legacy Informatica and new AWS Glue to validate data parity before full production cutover.

- Conducted foundational PoC establishing migration feasibility, technical architecture, and risk assessment.
- Built data transfer and validation pipelines using Informatica PowerCenter and SQL Server for high-traffic Insurance customer data.

Tata Consultancy Services (TCS)

Mar 2016 – Jun 2022

IT Analyst

- Delivered enterprise data warehousing solutions using Informatica PowerCenter, Snowflake, and SAP BOXI reporting across multi-year engagements.
- Led migration from Informatica PowerCenter to Informatica Cloud (IICS), driving modernization of legacy ETL infrastructure.
- Developed and maintained ETL workflows processing large-scale structured data for downstream warehousing and reporting use cases.

PORTFOLIO PROJECT

DataForge — Conversational BI Assistant

Coco Quest · Snowflake Northstar Event | Jul 2026

- Built a natural language BI assistant using Streamlit-in-Snowflake, enabling business users to query data via plain English without writing SQL.
- Implemented Cortex-powered query routing to semantic views or LLM-generated SQL dynamically across SALES, HR, and FINANCE schemas.
- Engineered production-grade guardrails: SQL injection mitigation, prompt injection defense, row-limit controls, schema context isolation, and confidence scoring.

GitHub: github.com/ramakrushna1994/Snowflake-BI-Assistant

CORE SKILLS

Platforms	Databricks, Snowflake, AWS Glue, ADLS Gen2, Informatica Cloud (IICS), Informatica PowerCenter
Languages	PySpark, Python, SQL
Streaming	Apache Kafka
Storage	Delta Lake, Medallion Architecture
Integration	REST API, XML Parsing, JSON Processing, Async Callback Patterns
Security	PII Tokenization, AES-256 Encryption, Azure Key Vault, Field-Level Encryption / Decryption
DevOps	Databricks Asset Bundles (DAB), CI/CD, GitHub Actions
Legacy	Oracle, Unix Shell Scripting, SAP BOXI

CERTIFICATIONS

- DP-750: Microsoft Azure Databricks Solution Expert (*In Progress*)
- Microsoft Azure Fundamentals (AZ-900)

EDUCATION

B.Tech / B.E.

Biju Patnaik University of Technology (BPUT) | 2015