

Ramakrushna Dash

Senior Data Engineer

ramakrushna.work@gmail.com | +91 9776955226 | linkedin.com/in/ramakrushna-dash-de | Bhubaneswar, India

PROFESSIONAL SUMMARY

Senior Data Engineer with ~10 years of experience architecting cloud-native data platforms across Insurance and Financial Services domains. Specializes in real-time Kafka streaming pipelines, lakehouse architecture on Databricks and Snowflake, and enterprise-grade PII tokenization and encryption. Proven track record modernizing legacy Informatica ETL ecosystems to scalable, event-driven solutions on Azure and AWS.

WORK EXPERIENCE

Chubb — Insurance Domain

Senior Data Engineer | Nov 2024 – Present

Phase 2 — Modern Streaming Platform (Mar 2025 – Present)

- Architected and delivered 3 production-grade streaming pipelines ingesting real-time JSON events from Apache Kafka into Databricks, processed via medallion architecture, and loaded to Snowflake.
- Engineered ChubbIO integration — consumed Kafka JSON event streams into Databricks bronze layer using PySpark and Delta Lake on ADLS Gen2.
- Developed DuckCreek integration using an async callback API pattern: receives JSON payload with async ID, triggers DuckCreek callback API to retrieve XML policy documents, parses XML, and loads structured data to Snowflake via Databricks.
- Built SQL Server to Snowflake migration pipeline in Databricks, handling embedded XML parsing across 2 columns before bronze layer ingestion.
- Implemented Chubb-standard PII tokenization and field-level encryption/decryption across all PII-sensitive pipelines to meet enterprise compliance requirements.
- Managed CI/CD deployment and environment promotion using Databricks Asset Bundles (DAB).

Phase 1 — UW Reporting (Nov 2024 – Mar 2025)

- Built data marts integrating Mainframe systems, flat files, and JSON sources using Informatica Cloud (IICS) and Snowflake.
- Delivered foundational data layer consumed by Power BI dashboards supporting underwriting business decisions and KPI reporting.
- Automated ETL pipelines using Python scripting, improving pipeline reliability and reducing manual intervention.

LTIMindtree

Specialist Data Engineer | Dec 2022 – Oct 2024

- Led end-to-end migration of 50+ Informatica PowerCenter workflows to AWS Glue for a large Insurance client, modernizing legacy ETL infrastructure to cloud-native architecture.
- Migrated initial workflow subset to production; implemented parallel load reconciliation — running dual pipelines across legacy Informatica and new AWS Glue to validate data parity before full cutover.
- Conducted foundational PoC establishing migration feasibility, technical architecture, and risk assessment.
- Built data transfer and validation pipelines using Informatica PowerCenter and SQL Server for high-traffic Insurance customer data.

Tata Consultancy Services (TCS)

IT Analyst | Mar 2016 – Jun 2022

- Delivered enterprise data warehousing solutions using Informatica PowerCenter, Snowflake, and SAP BOXI reporting across multi-year engagements.
- Led migration from Informatica PowerCenter to Informatica Cloud (IICS), driving modernization of legacy ETL infrastructure.
- Developed and maintained ETL workflows processing large-scale structured data for downstream warehousing and reporting use cases.

PORTFOLIO PROJECT

Snowflake Conversational BI Assistant — Coco Quest

Personal Project — Competitive Coding Contest Entry | Jul 2026

- Built a natural language BI assistant using Streamlit-in-Snowflake, enabling business users to query data via plain English without writing SQL.
- Implemented Cortex-powered query routing to semantic views or LLM-generated SQL dynamically across SALES, HR, and FINANCE schemas.
- Engineered production-grade guardrails: SQL injection mitigation, prompt injection defense, row-limit controls, schema context isolation, and conversation memory.

GitHub: github.com/ramakrushna1994/Snowflake-BI-Assistant

CORE SKILLS

Platforms: Databricks, Snowflake, AWS Glue, ADLS Gen2, Informatica Cloud (IICS), Informatica PowerCenter

Languages: PySpark, Python, SQL

Streaming: Apache Kafka

Storage: Delta Lake, Medallion Architecture

Integration: REST API, XML Parsing, JSON Processing, Async Callback Patterns

Security: PII Tokenization, Field-Level Encryption / Decryption

DevOps: Databricks Asset Bundles (DAB), CI/CD

Legacy: Oracle, Unix Shell Scripting, SAP BOXI

CERTIFICATIONS

- DP-750: Microsoft Azure Databricks Solution Expert (*In Progress*)
- Microsoft Azure Fundamentals (AZ-900)

EDUCATION

B.Tech / B.E.

Biju Patnaik University of Technology (BPUT) | 2015